

Advanced Topics

Chapter-16

The Java Memory Model

Upcode Software
Engineer Team

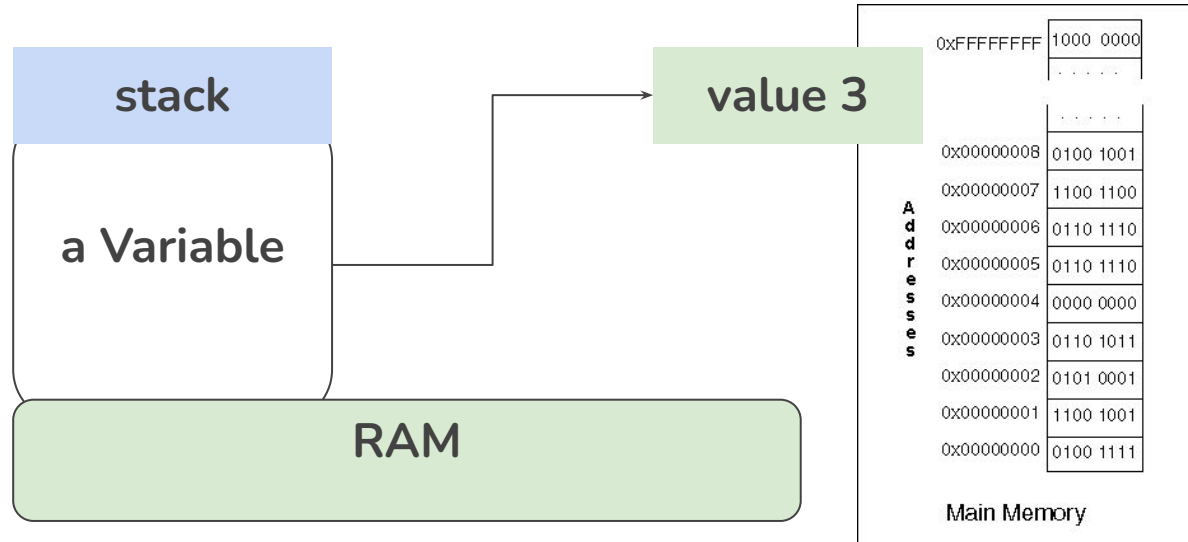


CONTENT

1. What is memory model ?
2. Why would I want one ?
3. Safe publication
4. Solution
5. Reference

What is memory model (1/n)

A memory model addresses the question “Under what conditions does a thread that reads a Variable see the value 3?”




What is memory model (2/n)

- **The main memory (or simply the memory)** is where variables and other information are stored while a program runs.
- From the perspective of a program, **the computer's memory is a collection of bytes, each with an integer address**. For example, there is a byte with address 1, another with address 2, etc., up to a very large number.
- A program can **fetch the current contents of the byte at a given memory address** and it can store a given value into that byte.

address		data
0x8000	→	0xD5
0x8001	→	0xF3
0x8002	→	0x64
0x8003	→	0xA7

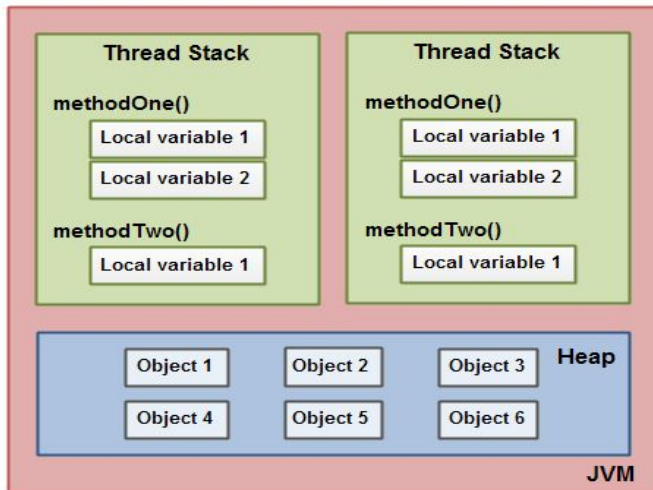
What is memory model (3/n)

- the JVM to maintain **within-thread as-if-serial semantics**: as long as the program has the same result.
- The JMM specifies the minimal guarantees **the JVM must make about when writes to variables become visible to other threads**. It was designed to **balance the need for predictability and ease of program development** with the realities of implementing high-performance JVMs on a wide range of popular processor architectures.

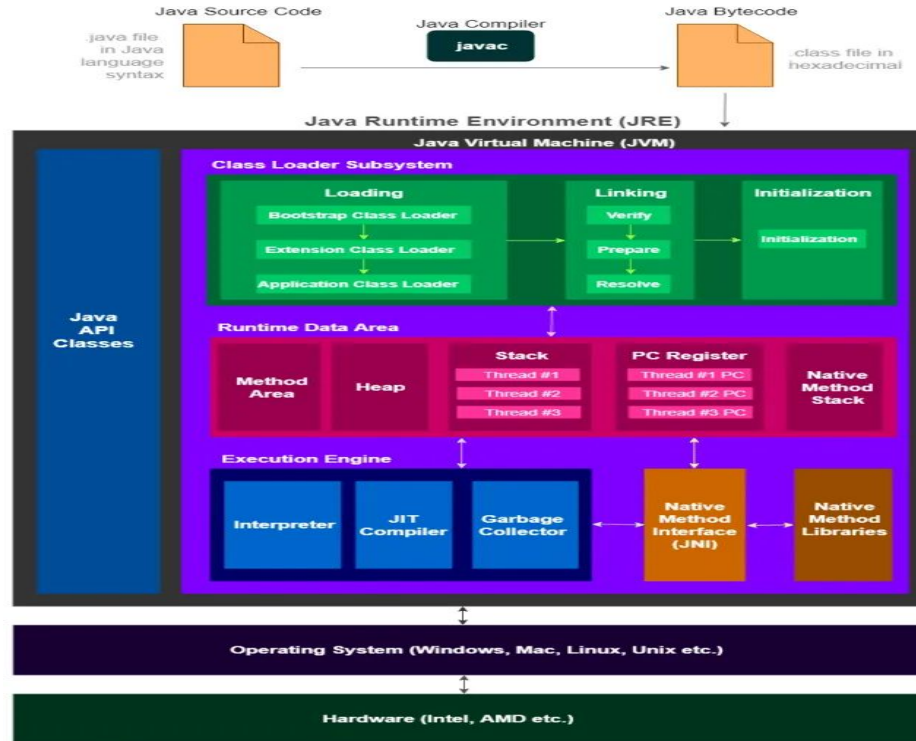


What is memory model (4/n)

- Some aspects of the JMM may be disturbing at first
- if you are not familiar with the tricks used by modern processors and compilers to squeeze extra performance out of your program.




What is memory model (5/n)



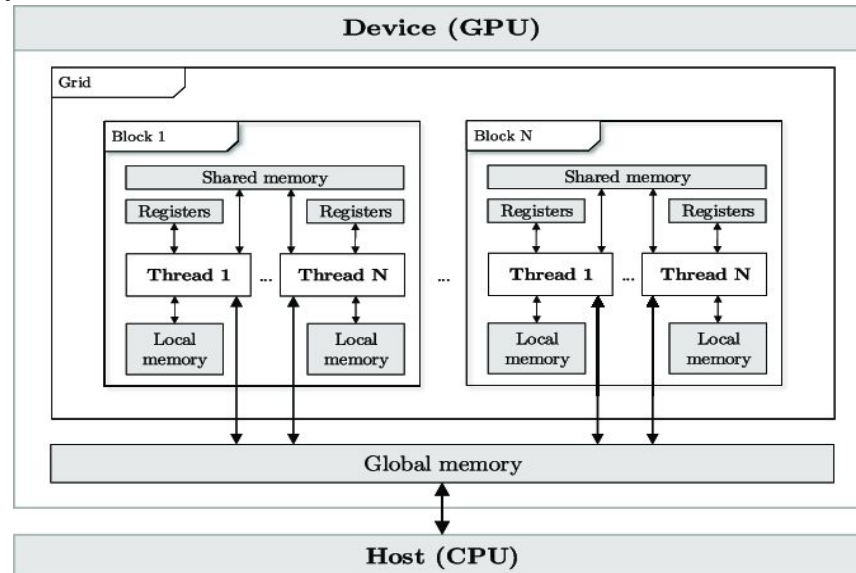
What is memory model (6/n)

- The Java Memory Model (JMM) defines the allowable behavior of multithreaded programs
- therefore describes when such reorderings are possible.
- It places **execution-time constraints** on the relationship between threads and main memory in order to achieve consistent and reliable Java applications.



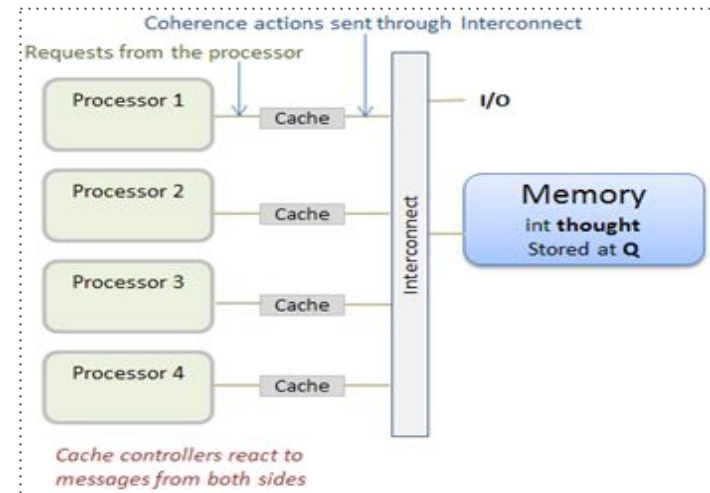
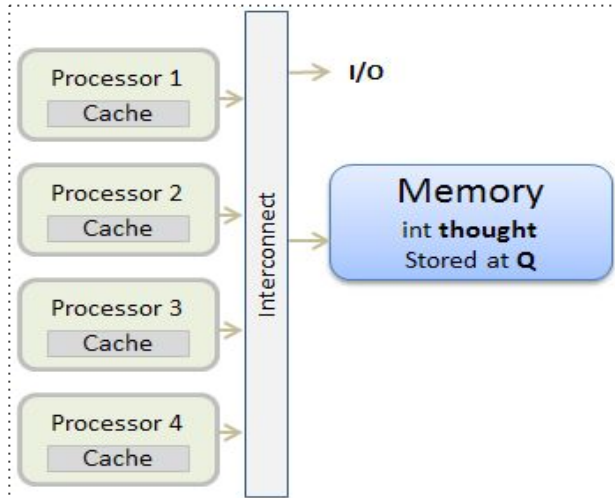
What is memory model (7/n) -Platform memory

- You must have used some of the following **JVM memory configurations** when running resource-intensive Java programs.
- Processor architectures provide varying degrees of *cache coherence*; some provide minimal guarantees that allow different processors to see different values for **the same memory location** at virtually any time.



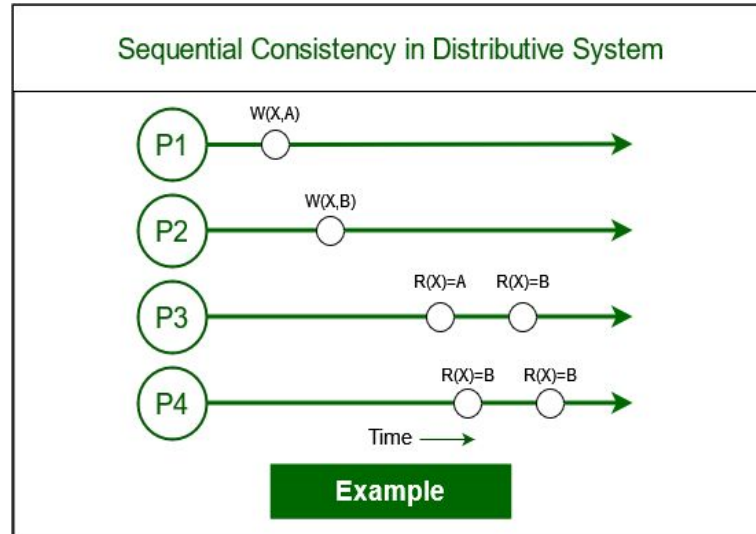
What is memory model (8/n) -Platform memory

- Most of the time this information is not needed, so processors relax their **memory-coherency guarantees to improve performance.**
- In order to shield **the Java developer from the differences between memory models across architectures**, Java provides its own memory model, and the JVM deals with **the differences between the JMM and the underlying platform's memory model by inserting memory barriers at the appropriate places.**




Sequential consistency

- Sequential consistency is a **conservative memory model that does not allow any instruction reordering on each core.**
- This prevents **many optimizations and degrades performance.** However, not all memory instructions on a single core need to preserve their program order



Sequential consistency

- **Sequential consistency** is a conservative memory model that does not allow any instruction reordering on each core.
- This prevents many optimizations and degrades performance.





The rules for happens-before

- **Program order rule.** Each action in a thread happens-before every action in that thread that comes later in the program order.
- **Monitor lock rule.** An unlock on a monitor lock happens-before every subsequent lock on that same monitor lock.
- **Volatile variable rule.** A write to a volatile field happens-before every subsequent read of that same field.
- **Thread start rule.** A call to `Thread.start` on a thread happens-before every action in the started thread.



The rules for happens-before

- **Thread termination rule.** Any action in a thread happens-before any other thread detects that thread has terminated, either by successfully return from `Thread.join` or by `Thread.isAlive` returning false.
- **Interruption rule.** A thread calling `interrupt` on another thread happens-before the interrupted thread detects the interrupt (either by having `InterruptedException` thrown, or invoking `isInterrupted` or `interrupted`).
- **Finalizer rule.** The end of a constructor for an object happens-before the start of the finalizer for that object.
- **Transitivity.** If A happens-before B, and B happens-before C, then A happens-before C

The rules for happens-before

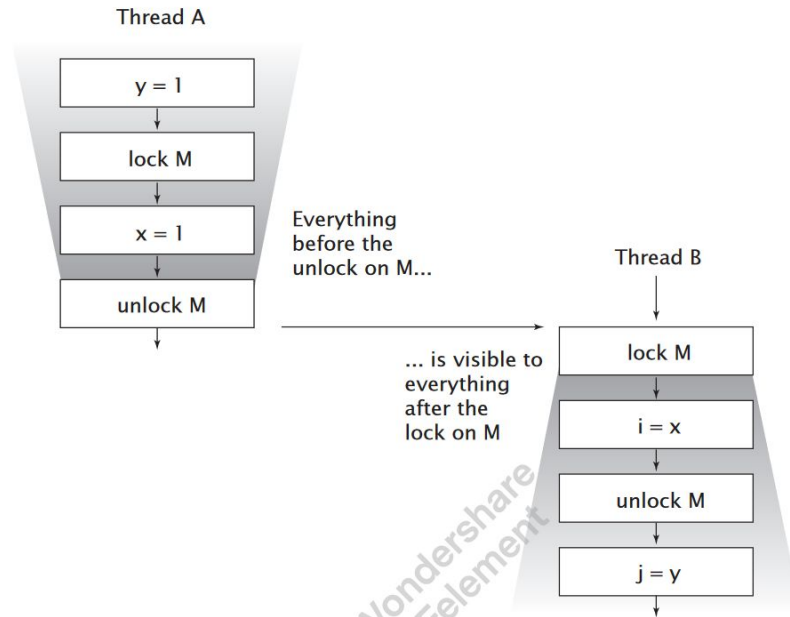


FIGURE 16.2. Illustration of *happens-before* in the Java Memory Model.



Piggybacking on synchronization



Summary

- the most common reasons to use **threads** is to exploit **multiple processors**, in

Resources





Reference

1. Java Concurrency book.
2. <https://medium.com/platform-engineer/understanding-java-memory-model-1d0863f6d973>
3. <https://medium.com/@jojooooo/exploring-a-base-spring-boot-application-with-java-21-virtual-thread-spring-security-flyway-c0fde13c1eca>
4. <https://jenkov.com/tutorials/java-concurrency/java-memory-model.html>



Thank you!

Presented by

Hamdamboy Urunov

(hamdamboy.urunov@gmail.com)